

Dhruv Susheelkar

(510) 579-8139 | dsusheelkar@ucsd.edu | [linkedin.com/in/dhruv-susheelkar](https://www.linkedin.com/in/dhruv-susheelkar) | github.com/3DSA | U.S. Citizen

EDUCATION

University of California, San Diego 

B.S. in Computer Science, June 2025

SKILLS

Languages: Python, Go, Java, C++, C, SQL, Bash, JavaScript, TypeScript, HTML/CSS

ML/AI: LangGraph, LangChain, TensorFlow, PyTorch, Scikit-learn, Pandas, NumPy, Matplotlib, Jupyter, MCP, Agentic AI, Multi-Agent Systems, Prompt Engineering, Vector DBs

Frameworks: Flask, FastAPI, Streamlit, React

Cloud/DevOps: AWS (EC2, S3, Lambda, EKS, ECS), Docker, Kubernetes, Terraform, Ansible, Jenkins, Git, Linux, Serverless, CI/CD

Big Data & Monitoring: Apache Spark, Hadoop, Kafka, Airflow, Prometheus, Grafana, Tableau, Power BI

Databases/APIs: PostgreSQL, MySQL, MongoDB, REST, GraphQL, Postman

Tools/Practices: Agile/Scrum, TDD, Integration Testing, Jira, Confluence, HTTP/HTTPS, WebSockets

EXPERIENCE

A.I./M.L. Engineer

April 2025 – Present

Advantis @ Apple

Cupertino, CA

- Cut **\$8M** in annual cloud costs by engineering an AWS optimization script that identified and reallocated underutilized EC2 GPU workloads across ML teams.
- Reduced S3 storage expenses by **\$4M** through detailed usage analysis and automated cross tier migration.
- Accelerated root-cause analysis by **90%** and improved debugging precision by enhancing Apple's Agentic AI system with adaptive iteration control, severity-aware scoring, and intelligent log retrieval.
- Developed an internal ML Agent leveraging Apple's MCP framework, accelerating engineering workflows and reducing dev cycle time
- Achieved real time and historical visibility into GPU resource usage by developing a Lambda based data pipeline that ingested cross account metrics into a PostgreSQL time-series database.
- Saved **\$3M** annually by migrating observability stack to Apple native tools, eliminating vendor lock in and improving precision.
- Replaced reliance on third party monitoring tools and improved telemetry accuracy by designing and deploying Go based collectors as node agents that streamed system and GPU metrics into Apple's internal observability platform.

A.I Researcher

September 2025 – Present

HandshakeAI @ OpenAI

Remote

- Evaluated LLM outputs using rubric-driven frameworks to assess factuality, linguistic precision, & task alignment across web domains.
- Accelerated model improvement by identifying recurring failure patterns & translating them into research insights & protocol updates.

Software Engineer Intern

June 2024 – September 2024

Calix

Hybrid

- Streamlined debugging and accelerated development workflows by 30% through engineering a component versioning module that surfaced code regressions and mismatches early in the CI pipeline.
- Reduced integration errors by 40% by building a command-line tool to validate schema compliance before deployment.
- Improved CI pipeline responsiveness and cut dev turnaround time by 25% by automating Bamboo bug report routing.

A.I./M.L. Engineer Intern

February 2024 – June 2024

FlowAI

Remote

- Improved the decision making performance of adaptive systems by integrating Reinforcement Learning algorithms, enabling dynamic strategy adjustment based on real-time feedback.
- Reduced lead data retrieval latency by 50% by developing a backend system optimized for concurrent processing.
- Boosted user contact automation by 35% through designing and deploying a CRM backend that unified data from APIs.

PROJECTS

Agentic Platform

Langgraph, Langchain, MCP, React, TypeScript, RDS Database, Python

- Enabled non-technical users to launch AI agents via natural language, cutting setup time from hours to minutes through automated prompt, schema, and workflow generation
- Drove enterprise adoption by automating agent provisioning, secure context retrieval, and multi-surface publishing (chat, APIs, embeds, share links).
- Built a scalable multi-agent architecture with LangGraph, LangChain, adaptive memory, and resilient integration patterns (OAuth2, retries, observability).
- Increased deployment success by 30% with robust monitoring and fallback mechanisms, ensuring SLA compliance.

Fracture Detection in X-ray Imaging

Python, TensorFlow, OpenCV, Vision Transformers

- Enabled automated fracture detection by building a deep learning model combining CNNs and Vision Transformers.
- Increased fracture classification accuracy from 85% to 92% by optimizing model architecture, directly improving diagnostic reliability.
- Improved spatial feature extraction by integrating EfficientNet B0, boosting model performance on the FracAtlas dataset.